


Message Passing and Node Classification

Outline

- **Main question today:** Given a network with labels on some nodes, how do we assign labels to all other nodes in the network?
- **Example:** In a network, some nodes are fraudsters and some other nodes are fully trusted. **How do you find the other fraudsters and trustworthy nodes?**
- We already discussed node embeddings as a method to solve this

Example: Node Classification




- Given labels of some nodes
- Let's predict labels of unlabeled nodes
- This is called semi-supervised node classification

Outline

- Today we will discuss some intuitions behind the framework: **message passing**
- Intuition: **Correlations** exist in networks.
 - In other words: Similar nodes are connected
 - Key concept is **collective classification**: Idea of assigning labels to all nodes in a network together
- We will look at some techniques today:
 - **Relational classification**
 - **Iterative classification**

Correlations in Networks


- Individual behaviors are **correlated** in the network
- **Correlation**: nearby nodes have the same color (belonging to the same class)




Correlations in Networks

- **Main types of dependencies that lead to correlation:**

Homophily




Influence



Homophily

- **Homophily**: The tendency of individuals to associate and bond with similar others
 - *“Birds of a feather flock together”*
 - It has been observed in a vast array of network studies, based on a variety of attributes (e.g., age, gender, organizational role, etc.)
 - **Example**: Researchers who focus on the same research area are more likely to establish a connection (meeting at conferences, interacting in academic talks, etc.)


Homophily



Homophily: Example

Example of homophily

- Online social network
 - Nodes = people
 - Edges = friendship
 - Node color = interests (sports, arts, etc.)
- People with the same interest are more closely connected due to homophily




(Easley and Kleinberg, 2010)

Influence


- **Influence:** Social connections can influence the individual characteristics of a person.
- **Example:** I recommend my musical preferences to my friends, until one of them grows to like my same favorite genres!

Influence



Classification with Networks

- How do we **leverage this correlation** observed in networks to help predict node labels?



How do we predict the labels for the nodes in grey?


Motivation

- **Similar nodes are typically close together or directly connected in the network:**
 - **Guilt-by-association:** If I am connected to a node with label X , then I am likely to have label X as well.
 - **Example: Malicious/benign web page:**
Malicious web pages link to one another to increase visibility, look credible, and rank higher in search engines

Motivation

- **Classification label** of a node v in network may depend on:
 - **Features** of v
 - **Labels** of the nodes in v 's neighborhood
 - **Features** of the nodes in v 's neighborhood

Semi-supervised Learning



Formal setting

Given:

- Graph
- Few labelled nodes

Find: class (**red**/**green**)
of remaining nodes

Main assumption:

There is homophily in
the network

Semi-supervised Learning

Example task:

- Let A be a $n \times n$ adjacency matrix over n nodes
- Let $Y = \{0, 1\}^n$ be a vector of **labels**:
 - $Y_v = 1$ belongs to **Class 1**
 - $Y_v = 0$ belongs to **Class 0**
 - There are **unlabeled** node needs to be classified
- **Goal**: Predict which **unlabeled** nodes are likely **Class 1**, and which are likely **Class 0**

Collective Classification

- **Many applications:**
 - Document classification
 - Part of speech tagging
 - Link prediction
 - Optical character recognition
 - Image/3D data segmentation
 - Entity resolution in sensor networks
 - Spam and fraud detection

Collective Classification

- **Intuition:** Simultaneous classification of interlinked nodes using correlations
- Probabilistic framework
- **Markov Assumption:** *the label Y_v of one node v depends on the labels of its neighbors N_v*

$$P(Y_v) = P(Y_v | N_v)$$

- Collective classification involves 3 steps:

Local Classifier	Relational Classifier	Collective Inference
<ul style="list-style-type: none">• Assign initial labels	<ul style="list-style-type: none">• Capture correlations between nodes	<ul style="list-style-type: none">• Propagate correlations through network

Collective Classification

Local Classifier

- Assign initial labels

Local Classifier: Used for initial label assignment

- Predicts label based on node attributes/features
- Standard classification task
- Does not use network information

Relational Classifier

- Capture correlations between nodes

Relational Classifier: Capture correlations

- Learns a classifier to label one node based on the labels and/or attributes of its neighbors
- This is where network information is used

Collective Inference

- Propagate correlations through network


Collective Inference: Propagate the correlation

- Apply relational classifier to each node iteratively
- Iterate until the inconsistency between neighboring labels is minimized
- Network structure affects the final prediction

Problem Setting

- How to predict the labels Y_v for the unlabeled nodes v (in grey color)?
- Each node v has a feature vector f_v
- Labels for some nodes are given (1 for green, 0 for red)
- **Task:** Find $P(Y_v)$ given all features and the network

$$P(Y_v) = ?$$



What next?

- We focus on semi-supervised node classification
- Intuition is based on **homophily**: Similar nodes are typically close together or directly connected
- **Techniques we will introduce:**
 - **Relational classification**
 - **Iterative classification**

Relation Classification and Iterative Classification

Collective Classification Models

- **Relational classifiers**
- Iterative classification

Probabilistic Relational Classifier

- **Basic idea:** Class probability Y_v of node v is a weighted average of class probabilities of its neighbors
- For **labeled nodes** v , initialize label Y_v with ground-truth label Y_v^*
- For **unlabeled nodes**, initialize $Y_v = 0.5$
- **Update** all nodes in a random order until convergence or until maximum number of iterations is reached

Probabilistic Relational Classifier

- **Update** for each node v and label c (e.g. 0 or 1)


$$P(Y_v = c) = \frac{1}{\sum_{(v,u) \in E} A_{v,u}} \sum_{(v,u) \in E} A_{v,u} P(Y_u = c)$$

- If edges have strength/weight information, $A_{v,u}$ can be the edge weight between v and u
 - $P(Y_v = c)$ is the probability of node v having label c
- **Challenges:**
 - Convergence is not guaranteed
 - Model cannot use node feature information

Initialization

Initialization:

- All labeled nodes with their labels
- All unlabeled nodes 0.5 (belonging to class 1 with probability 0.5) Let $P_{Y_1} = P(Y_1 = 1)$




1st Iteration, Update Node 3

- Update for the 1st Iteration:


- For node 3, $N_3 = \{1, 2, 4\}$

$$P(Y_3) = (0 + 0 + 0.5)/3 = 0.17$$



1st Iteration, Update Node 4

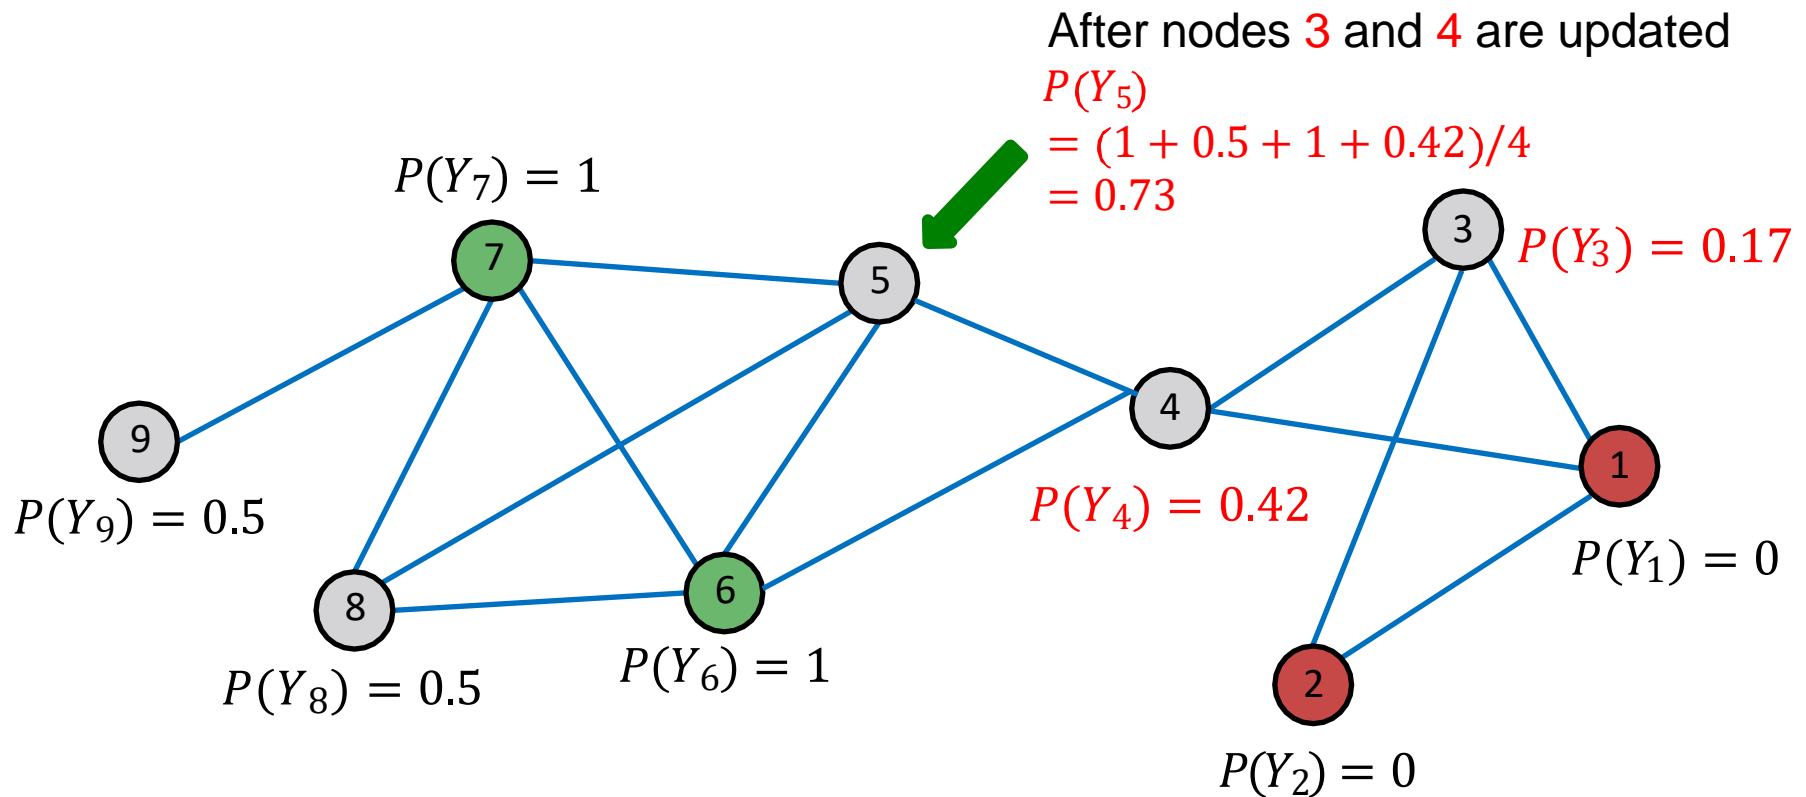
- Update for the 1st Iteration:
 - For node 4, $N_4 = \{1, 3, 5, 6\}$



After Node 3 is updated


1st Iteration, Update Node 5

- Update for the 1st Iteration:
 - For node 5, $N_5 = \{4, 6, 7, 8\}$




End of 1st Iteration

- After Iteration 1
(a round of updates for all unlabeled nodes)




After 2nd Iteration

- After Iteration 2




After 3rd Iteration

- After Iteration 3




After 4th Iteration

- After Iteration 4



Convergence

- All scores stabilize after 4 iterations. We therefore predict:
 - **Nodes 4, 5, 8, 9 belong to class 1** ($P_{Y_v} > 0.5$)
 - **Nodes 3 belong to class 0** ($P_{Y_v} < 0.5$)



Collective Classification Models

- Relational classifiers
- **Iterative classification**

Iterative Classification

- Relational classifiers **do not use node attributes**. How can one leverage them?
- **Main idea of iterative classification:** Classify node v based on its **attributes f_v** as well as **labels z_v** of neighbor set N_v

Iterative Classification

- **Input: Graph**

- f_v : feature vector for node v
- Some nodes v are labeled with Y_v

- **Task:** Predict label of unlabelled nodes

- **Approach: Train two classifiers:**


- $\phi_1(f_v)$ = Predict node label based on node feature vector f_v
- $\phi_2(f_v, z_v)$ = Predict label based on node feature vector f_v and summary z_v of labels of v 's neighbors.

Computing Summary z_v

How do we compute the summary z_v of labels of v 's neighbors N_v ?

Ideas: z_v = **vector captures labels around node v**

- Histogram of the number (or fraction) of each label in N_v
- Most common label in N_v
- Number of different labels in N_v



Architecture of Iterative Classifiers

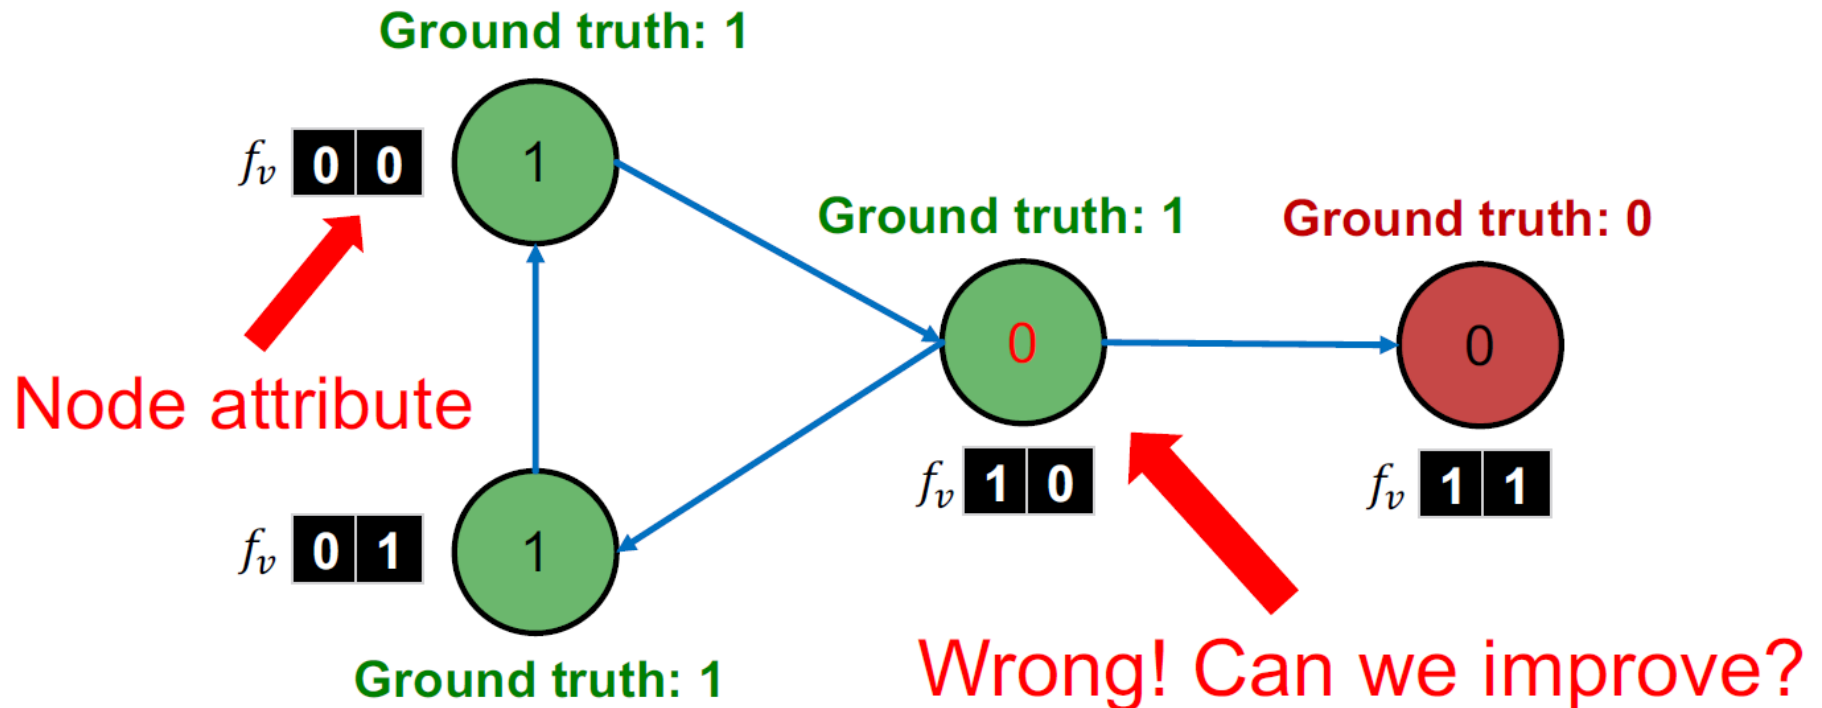
- **Phase 1: Classify based on node attributes alone**
 - On a **training set**, train classifier (e.g., linear classifier, neural networks, ...):
 - $\phi_1(f_v)$ to predict Y_v based on f_v
 - $\phi_2(f_v, z_v)$ to predict Y_v based on f_v and summary z_v of labels of v 's neighbors
- **Phase 2: Iterate till convergence**
 - On **test set**, set labels Y_v based on the classifier ϕ_1 , compute z_v and **predict the labels with ϕ_2**
 - **Repeat** for each node v
 - Update z_v based on Y_u for all $u \in N_v$
 - Update Y_v based on the new z_v (ϕ_2)
 - Iterate until class labels stabilize or max number of iterations is reached
 - Note: Convergence is not guaranteed

Example: Web Page Classification

- **Input:** Graph of web pages
- **Node:** Web page
- **Edge:** Hyper-link between web pages
 - **Directed edge:** a page points to another page
- **Node features:** Webpage description
 - For simplicity, we only consider 2 binary features
- **Task:** Predict the topic of the webpage

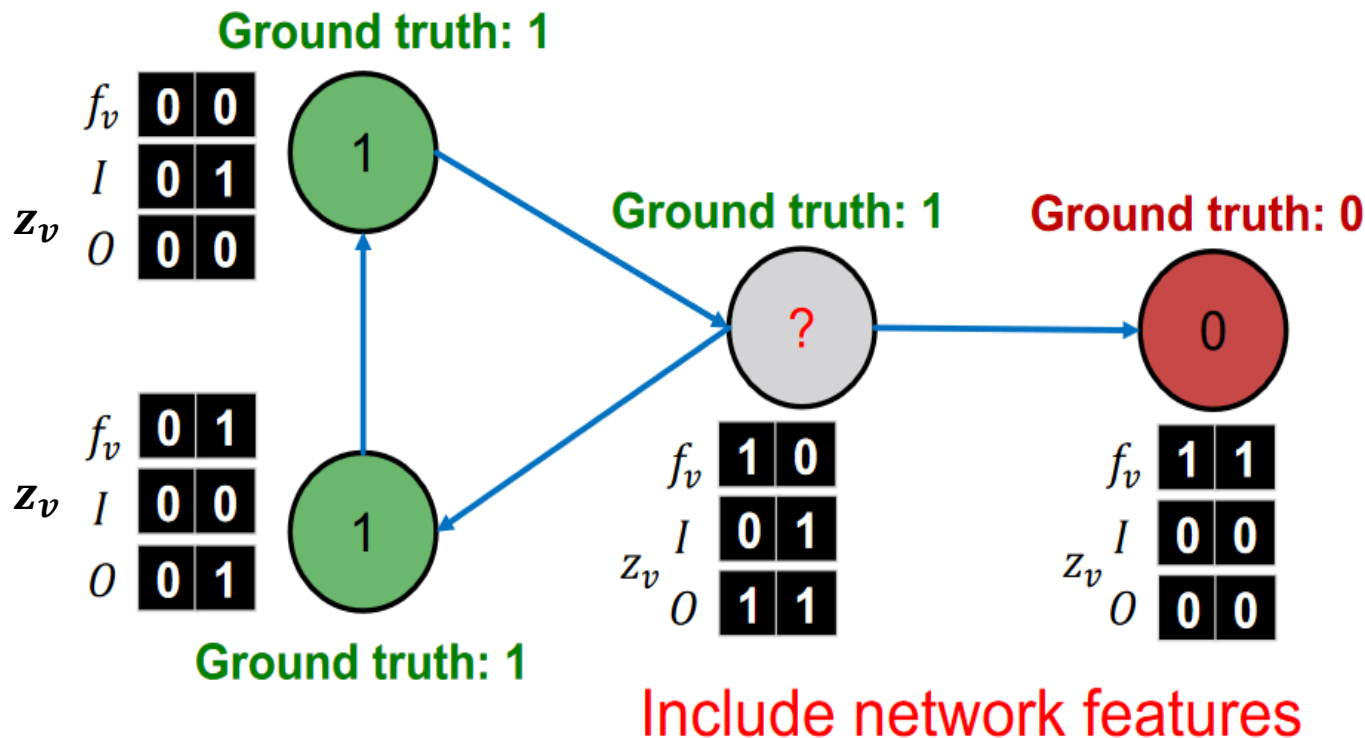
Example: Web Page Classification

- **Baseline:** train a classifier (e.g., linear classifier) to classify pages based on binary node attributes.



Example: Web Page Classification

- Each node maintains **vectors z_v of neighborhood labels**:
 - I = **Incoming** neighbor label information vector.
 - O = **Outgoing** neighbor label information vector.
 - $I_0 = 1$ if at least one of the incoming pages is labelled 0.
Similar definitions for I_1 , O_0 , and O_1



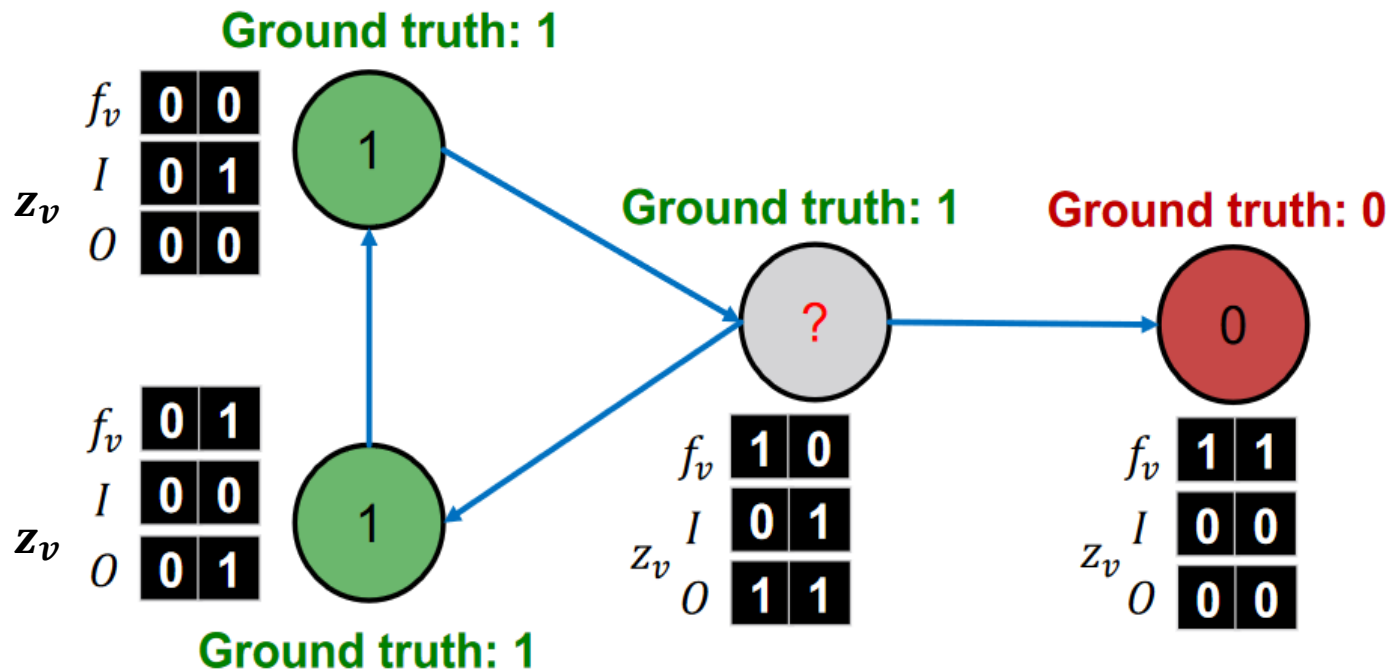
Iterative Classifier – Step 1

- On **training labels**, train two classifiers:
 - Node attribute vector only: $\phi_1(f_v)$
 - Node attribute and link vectors z_v : $\phi_2(f_v, z_v)$

1. Train classifiers

2. Apply classifier to unlab. set
3. Iterate

4. Update relational features z_v
5. Update label Y_v



Iterative Classifier – Step 2

- On the **unlabeled set**:
 - Use trained node feature vector classifier ϕ_1 to set Y_v


1. Train classifiers

2. Apply classifier to unlabeled set

3. Iterate

4. Update relational features z_v

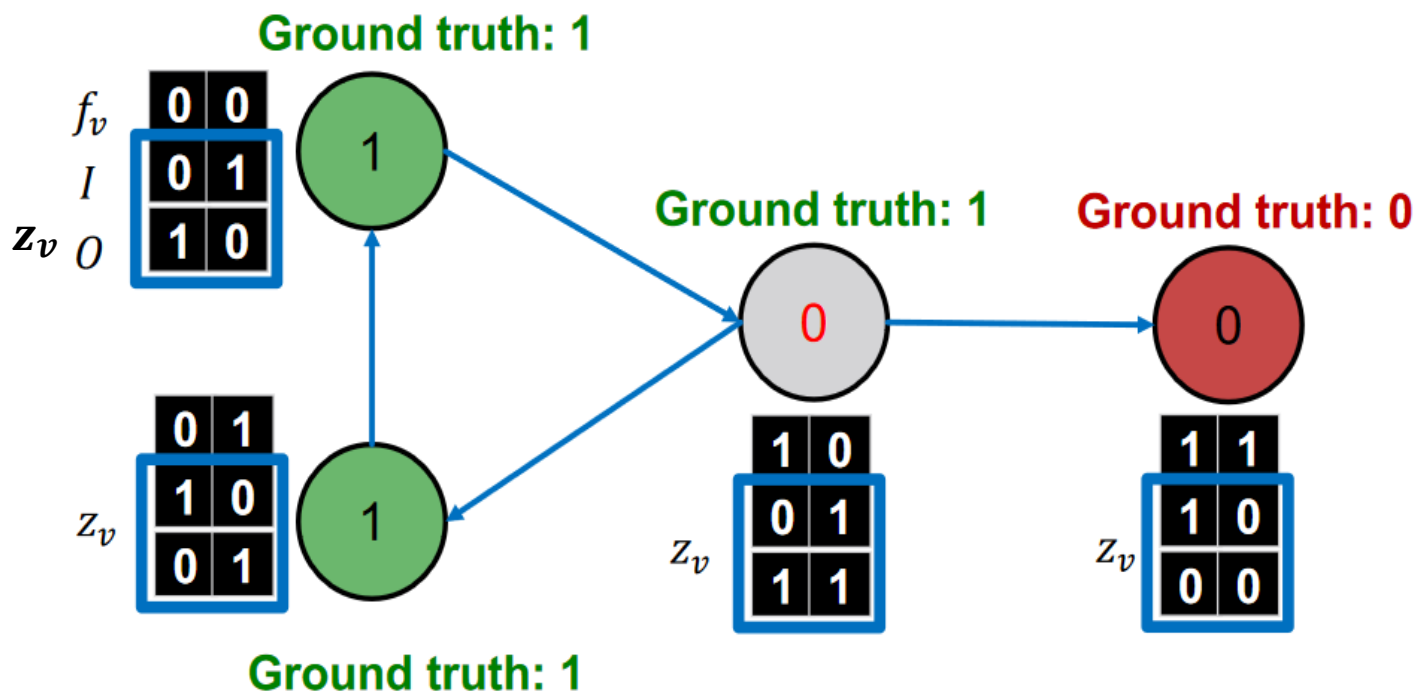
5. Update label Y_v



Iterative Classifier – Step 4

■ Update z_v for all nodes:

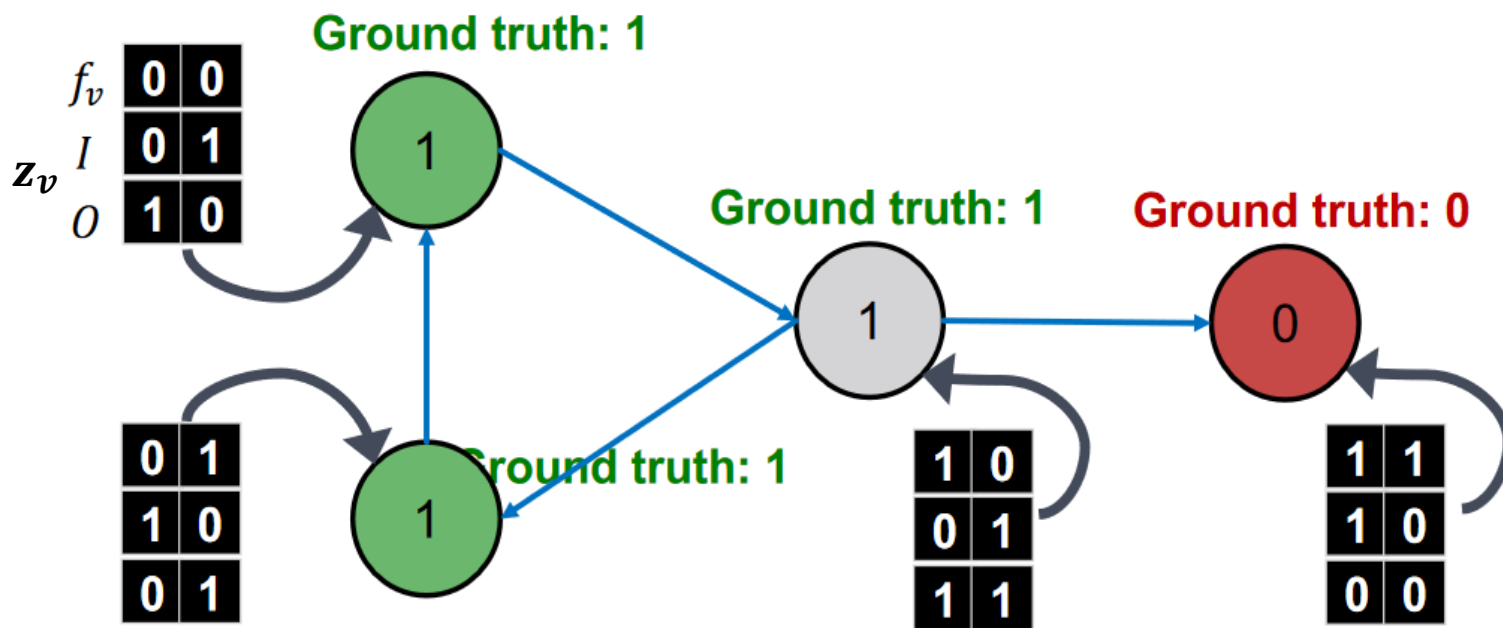
1. Train classifiers
2. Apply classifier to unlab. set
3. Iterate
 4. Update relational features z_{12}
 5. Update label Y_v



Iterative Classifier – Step 5

- Re-classify all nodes with ϕ_2 :

1. Train classifiers
2. Apply classifier to unlab. set
3. Iterate
 4. Update relational features z_v
 5. Update label y_v



Now it's correct prediction!

Iterative Classifier - Iterate

■ Continue until convergence

- Update z_v based on Y_v
- Update $Y_v = \phi_2(f_v, z_v)$

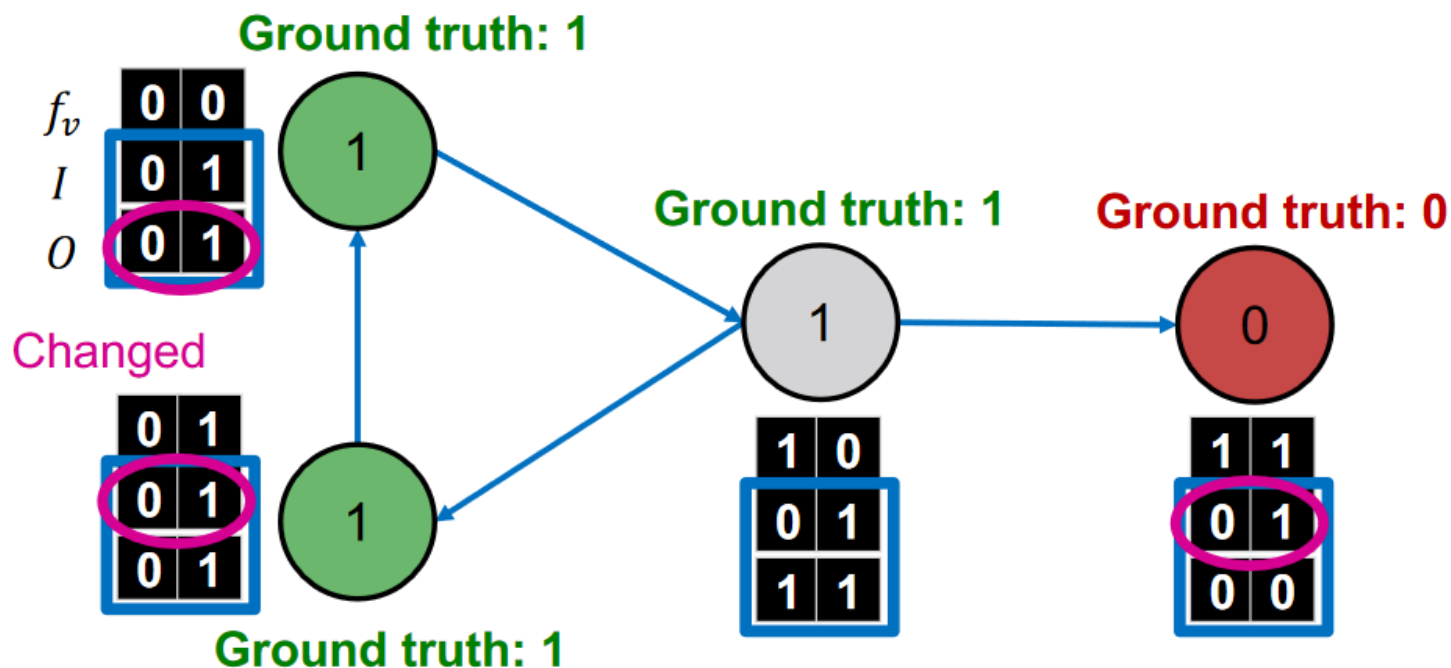
1. Train classifiers

2. Apply classifier to unlab. set

3. Iterate

4. Update relational features z_v


5. Update label Y_v



Iterative Classifier – Prediction

■ Stop iteration

- After convergence or when maximum iterations are reached



Summary

- We talked about 2 approaches to collective classification
- **Relational classification**
 - Iteratively update probabilities of node belonging to a label class based on its neighbors
- **Iterative classification**
 - Improve over collective classification to handle attribute/feature information
 - Classify node i based on its **features** as well as **labels** of neighbors